

Maximizing ORACLE® CLOUD Buffer Cache Throughput

Craig Shallahamer
craig@orapub.com



ORACLE®
ACE Director



ORAPUB



This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.



Top 5 Timed Foreground Events

Event	Waits	Time(s)	Avg wait (ms)	% DB time	Wait Class
DB CPU		8,344		56.01	
free buffer waits	326,788	4,689	14	31.48	Configuration
db file scattered read	57,041,171	1,065	0	7.15	User I/O
log buffer space	9,971	979	98	6.57	Configuration
log file switch (private strand flush incomplete)	1,014	391	386	2.63	Configuration



Background Wait Events

- ordered by wait time desc, waits desc (idle events last)
- Only events with Total Wait Time (s) >= .001 are shown
- %Timeouts: value of 0 indicates value was < .5%. Value of null is truly 0

Event	Waits	%Time-outs	Total Wait Time (s)	Avg wait (ms)	Waits /txn	% bg time
db file async I/O submit	9,211	0	4,200	456	6.12	51.86
log file parallel write	48,424	0	2,895	60	32.18	35.75
control file parallel write	11,982	0	747	62	7.96	9.22
db file single write	3,608	0	152	42	2.40	1.87
log file single write	1,576	0	50	32	1.05	0.62

About Me...

- Long time Oracle DBA
- Specialize in Oracle Database performance and predictive analysis
- Performance researcher
- Blogger: A Wider View About Oracle Performance Tuning
- Author: Oracle Performance Firefighting and Forecasting Oracle Performance.
- Conference speaker
- Teacher and mentor
- Oracle ACE Director
- IOUG DBA Track Manager

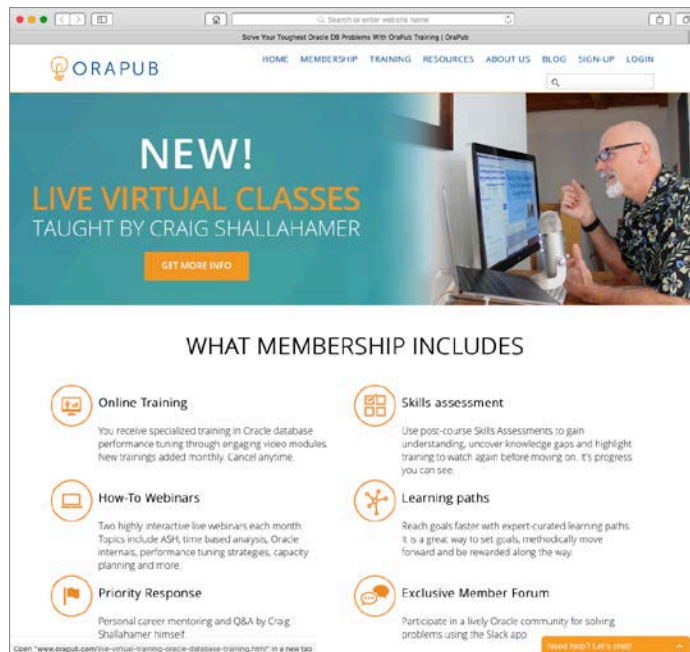


(c)OraPub, Inc.

ORAPUB

Maximizing BC Throughput

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.



(c)OraPub, Inc.

ORAPUB

Maximizing BC Throughput

OraPub works closely with DBAs who want to take their Oracle tuning skills to the next level.

So, you can solve the toughest Oracle DB problems.

How connect with Craig and OraPub



Community with a common goal



craig@orapub.com



[@OraPubInc](https://twitter.com/OraPubInc)



[OraPub.Com](https://orapub.com): Everything starts here!



Connect and network


This presentation was given by Craig Shallahamer (craig@orapub.com)
at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Key Topics



- Key buffer cache memory structures
- LRU list internals and algorithm
- LRU processing algorithm interesting observations
- How free buffer waits occur
- Free buffer wait solutions
- Going deeper

Key Buffer Cache Memory Structures



This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Free, Dirty, and Pinned...

- **Free/mirrored.** When a block's cached image mirrors the actual block on disk. The cached block (i.e., the buffer) can be replaced by another block and it can be referenced by any Oracle process. When the contents of a dirty buffer has been written to disk, it is once again a mirror of the actual block and placed back on the LRU end of its LRU list. An empty buffer is also called a free buffer. There are multiple buffer states (`x$bh.state`) related to a free/mirrored buffer.
- **Dirty.** When a block's cached image is not the same as its database file image. After a block has been changed (byte level) and has not yet been written back to its database file, it is deemed "dirty." (`v$bh.dirty='Y'`)
- **Pinned.** A buffer is pinned to keep it from being replaced. You don't want a block you are referencing to be replaced! Pinning can also be used to help ensure serial access to the buffer.

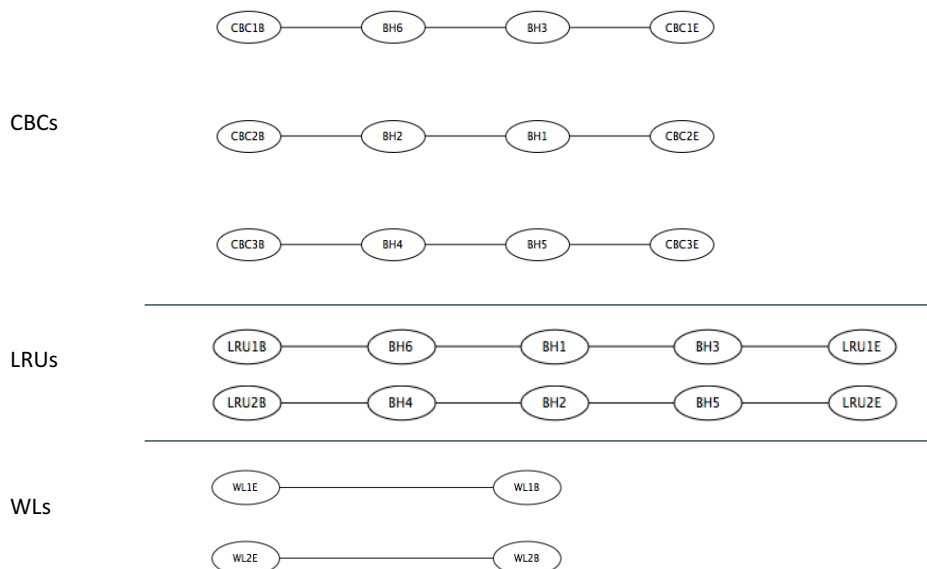
There are three key buffer cache lists.



- **Hash/Cache buffer chains** are used to quickly determine if a block is in the buffer cache.
- **LRU (LRU) chains** are used to keep popular buffers in the cache and to find free buffers.
- **Write (LRU-W) lists** are used by the DBWR to batch write dirty buffers. The write list is also called the *dirty list*.
- **Buffer headers** contain pointers to the physical block, the buffer cache location, and its position on its LRU, write, and hash chain list.

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Buffer headers are on many lists.



The buffer head view, v\$bh in 12c.


No changes in x\$bh. TCH column still only in x\$bh.

```
SQL> desc v$bh
```


Name	Null?	Type
FILE#		NUMBER
BLOCK#		NUMBER
CLASS#		NUMBER
STATUS		VARCHAR2 (10)
XNC		NUMBER
FORCED_READS		NUMBER
FORCED_WRITES		NUMBER
LOCK_ELEMENT_ADDR		RAW (8)
LOCK_ELEMENT_NAME		NUMBER
LOCK_ELEMENT_CLASS		NUMBER
DIRTY		VARCHAR2 (1)
TEMP		VARCHAR2 (1)
PING		VARCHAR2 (1)
STALE		VARCHAR2 (1)
DIRECT		VARCHAR2 (1)
NEW		CHAR (1)
OBJD		NUMBER
TS#		NUMBER
LOBID		NUMBER
CACHEHINT		NUMBER
FLASH_CACHE		VARCHAR2 (7)
CELL_FLASH_CACHE		VARCHAR2 (7)
CON_ID		NUMBER

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

The "real" stars of the buffer cache!




It's the lists and the buffer headers.

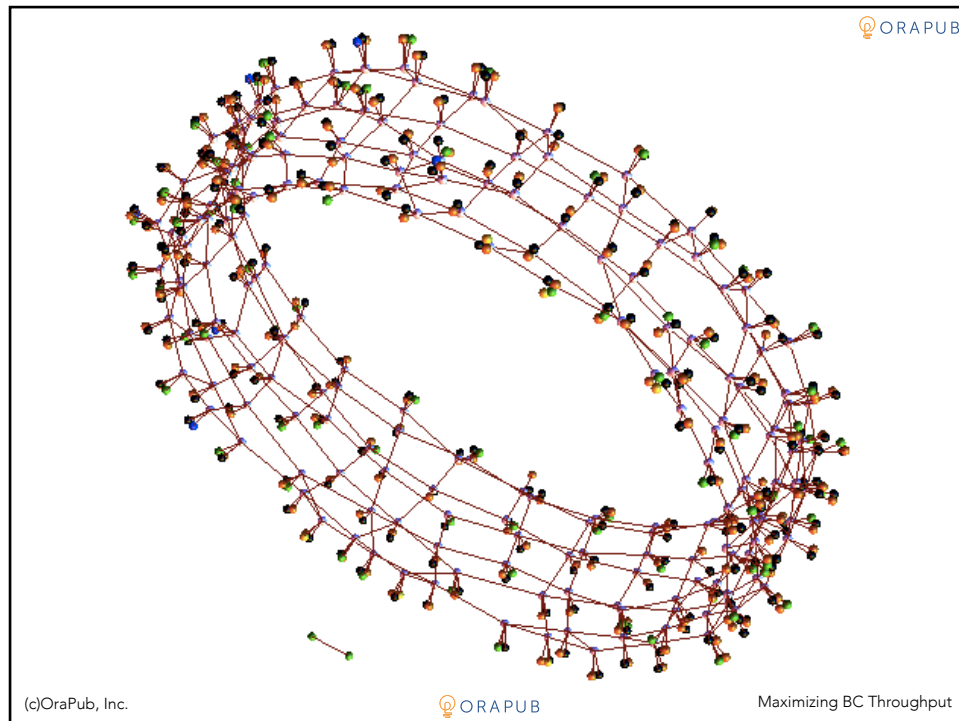


This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

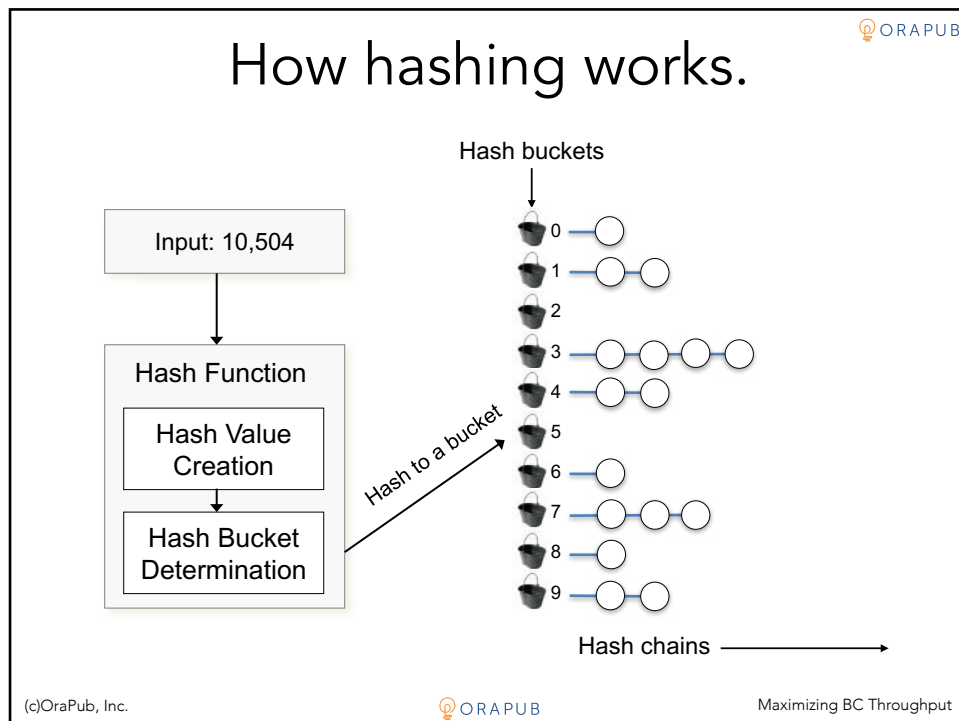
Visualizing 400 BH fully connected.




400 BHs
350 CBCs
2 LRUs
20% dirty
No cloning



This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.



Least Recently Used List Internals and Algorithm




This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Least recently used chains; LRUs.

- LRUs are used to cache hot buffers and to quickly supply free buffers for replacement to server processes.
- Cached buffers are divided into working sets.
- Each working set has an associated LRU chain or list.
- LRUs are made up of cached buffer headers ordered by *popularity*, descending. (not really...)
- Each LRU can contain free and dirty buffer headers. And they can be pinned or not pinned.
- Each DBWR is associated with one or more LRUs.
- There appears to be one LRU latch (`_db_block_lru_latches`) for each LRU chain (`x$kcwds`). A typical value is less than 100 LRUs.

The LRUs



```
SQL> select count(*) from x$kcwbds;
```

```
  COUNT(*)
-----
        48
```

```
SQL> @ipx %lru%latch%
```

Instance Parameter and Value	Description	Dflt?
<code>_db_block_lru_latches</code> = 48	number of lru latches	TRUE

o12.1

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Simulator LRU latch?

```
SQL> 1
    1 select name, count(*)
    2 from   v$latch_children
    3 where  name like '%lru%'
    4* group by name
SQL> /
```

NAME	COUNT (*)
simulator lru latch	48
cache buffers lru chain	96

2 rows selected.

o12.1


Are there 48 or 96 LRU latches?

Focus on cache buffers lru chain

```
SQL> def name
DEFINE NAME                = "cache buffers lru chain" (CHAR)
SQL> /
1  select name, addr, gets, sleeps
2  from    v$latch_children
3  where   name like '&name'
4* order by 1,2,4,3
SQL> /
```

NAME	ADDR	GETS	SLEEPS
...			
cache buffers lru chain	000000006D55CB20	0	0
cache buffers lru chain	000000006D55D1D0	2117	0
cache buffers lru chain	000000006D55D298	0	0
cache buffers lru chain	000000006D721018	58	0
cache buffers lru chain	000000006D7210E0	0	0
cache buffers lru chain	000000006D721790	58	0
cache buffers lru chain	000000006D721858	0	0
cache buffers lru chain	000000006D721F08	58	0
cache buffers lru chain	000000006D721FD0	0	0
cache buffers lru chain	000000006D722680	58	0

96 rows selected.



o12.1


This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Focus on used LRU chain latches

```
SQL> /
1  select name, addr, gets, sleeps
2  from    v$latch_children
3  where   name like '&name'
4    and  gets > 0
5* order by 1,2,4,3
/
```

NAME	ADDR	GETS	SLEEPS
...			
cache buffers lru chain	000000006CCDF8B0	58	0
cache buffers lru chain	000000006CCE0028	58	0
cache buffers lru chain	000000006CCE07A0	58	0
cache buffers lru chain	000000006CCE0F18	58	0
cache buffers lru chain	000000006CCE1690	58	0

48 rows selected.



o12.1

Let's summarize what we see

- We can see the number LRUs defined:
 - select count(*) from **x\$kcbwds**
- The number latches in use is set by the instance parameter `_db_block_lru_latches`
- Do not determine the number of LRU latches simply from `v$latch_children` entries.

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Check our assumptions on a big production system – 35GB BC

```
SQL> 1
      2 select name, count(*)
      3 from   v$latch_children
      4 where name like '%lru%'
      5* group by name
SQL> /
```

NAME	COUNT (*)
cache buffers lru chain	1536
simulator lru latch	768

2 rows selected.

o12.1

Are there 1536 or 768 LRU latches?

Focus on used simulator lru latches

```
SQL> def name
DEFINE NAME          = "simulator lru latch" (CHAR)
SQL> 1
  1 select name, addr, gets, sleeps
  2 from   v$latch_children
  3 where  name like '&name'
  4       and gets > 0
  5* order by 1,2,4,3
SQL> /
```

NAME	ADDR	GETS	SLEEPS
...			
simulator lru latch	0700010DA7C2C848	16	0
simulator lru latch	0700010DA7C2CF48	16	0
simulator lru latch	0700010DA7C2D648	16	0
simulator lru latch	0700010DA7C2DD48	16	0
simulator lru latch	0700010DA7C2E448	16	0
simulator lru latch	0700010DA7C2EB48	16	0
simulator lru latch	0700010DA7C2F248	16	0
simulator lru latch	0700010DA7C2F948	16	0
simulator lru latch	0700010DA7C30048	16	0

768 rows selected.

o12.1

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Focus on used LRU chain latches

```
SQL> def name
DEFINE NAME          = "cache buffers lru chain" (CHAR)
SQL> 1
  1 select name, addr, gets, sleeps
  2 from   v$latch_children
  3 where  name like '&name'
  4       and gets > 0
  5* order by 1,2,4,3
SQL> def name
DEFINE NAME          = "simulator lru latch" (CHAR)
```

NAME	ADDR	GETS	SLEEPS
...			
cache buffers lru chain	0700010DA7C2C9B0	1354	0
cache buffers lru chain	0700010DA7C2D0B0	1354	0
cache buffers lru chain	0700010DA7C2D7B0	1354	0
cache buffers lru chain	0700010DA7C2DEB0	1354	0
cache buffers lru chain	0700010DA7C2E5B0	1354	0
cache buffers lru chain	0700010DA7C2ECB0	1354	0
cache buffers lru chain	0700010DA7C2F3B0	1354	0
cache buffers lru chain	0700010DA7C2FAB0	1354	0
cache buffers lru chain	0700010DA7C301B0	1354	0

768 rows selected.

o12.1

What's the situation on your system?

```
col name format a25

select name, count(*)
from   v$latch_children
where  name like '%lru%'
group by name;

def name='cache buffers lru chain'
def name='simulator lru latch'


select name, addr, gets, sleeps
from   v$latch_children
where  name like '&name'
order by 1,2,4,3;

select name, addr, gets, sleeps
from   v$latch_children
where  name like '&name'
       and gets > 0
order by 1,2,4,3;
```

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.


Oracle's caching struggle...

- Standard least recently used algorithm
- Modified least recently used algorithm
- Touch-Count algorithm




The key TC components

- Midpoint insertion
- Touch count incrementation
- Buffer promotion
- Parameters for everything...



This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Demonstration



TC midpoint insertion details

- The buffer cache is divided into a *hot region* and a *cold region*.
- A *midpoint pointer* is maintained which moves to ensure proper regional block quantities.
- When a block is brought into the cache, it is placed in the “middle” of the LRU chain.
- A buffer naturally moves from the hot region into the cold region.

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Incrementing the touch count

- In concept, whenever a buffer is touched its TC is incremented.
- In reality, it does not work this way.
- No TC latch exist, thereby eliminating latch contention.
- To reduce rapid TC increases, the touch count will only be incremented once every 3 seconds.
- Buffer promotion and touch count increments are independent events.

Buffer promotion


- Keep in mind:
 - TC increment is buffer promotion independent.
 - Blocks are inserted at the midpoint.
- A buffer is promoted to the head of its current LRU when it's not pinned, its touch count is greater than 2 and:
 - A server process is looking for a free buffer -OR-
 - The DBWR process is looking for dirty buffers
- When buffer promotion does occur, the buffer's touch count is reset. This is important!

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Hot region to cold movement

- Regardless of its touch count,
- If a buffer crosses from the hot region into the cold region,
- It's touch count is reset to ONE.

LRU Processing Algorithm Interesting Observations



This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Does the 12c touch count still oscillate?

```
truncate table op_tch_results;
alter session set commit_wait=nowait;

declare
  loop_var number;
begin
  for loop_var in 1..1000
  loop
    insert into op_tch_results
      select loop_var, current_timestamp,
             dbarfil, dbablk, tch
      from   x$bhh
      where  dbarfil = 8
            and dbablk = 984
      ;
    commit;
    dbms_lock.sleep(1.0);
  end loop;
end;
/
```

The touch count instance parameters appear to be the same number, name, and default... no change.

But does the algorithm work the same?...

```
SQL> 1
      1* select current_ts ts, tch from op_tch_results order by 1;
SQL> /
```

TS	TCH
05.04.35.564918 PM	105
05.04.36.593226 PM	105
05.04.37.621009 PM	105
05.04.38.648773 PM	106
05.04.39.698656 PM	106
05.04.40.726701 PM	106
05.04.41.754430 PM	107
05.04.42.780681 PM	107
05.04.43.808713 PM	107
05.04.44.836228 PM	108
05.04.45.919045 PM	108
05.04.46.946001 PM	108
05.04.47.972823 PM	109
...	...
05.12.02.092964 PM	253
05.12.03.120867 PM	254
05.12.04.149798 PM	254
05.12.05.177229 PM	254
05.12.06.205295 PM	255
05.12.07.235103 PM	255
05.12.08.263965 PM	255
05.12.09.292339 PM	0
05.12.10.321628 PM	0
05.12.11.350899 PM	0
05.12.12.379346 PM	1
05.12.13.431346 PM	1
05.12.14.460496 PM	1
05.12.15.498567 PM	2

Note: This buffer was constantly being touched.

The workload was only LIQ, so there was never any need to find a non-popular free buffer.

Therefore, no buffer replacement occurred. And, no promotion therefore, the TCs were not reset to zero. Except... when it “rolled over.”

Notice:

1. The TC never exceeds 255.
2. The TC never increases more than once every three seconds.

o12.1

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

The TCH never exceeded 255

```
SQL> 1
      1 select dbarfil, dbablk, tch
      2 from   x$bh
      3* where  tch = ( select max(tch) from x$bh )
SQL> /
```

DBARFIL	DBABLK	TCH
6	1807676	255
6	1807299	255
6	1807532	255
9	31969	255
10	31915	255
9	31885	255
6	1808071	255
6	1807694	255
10	31941	255
6	19760	255
9	31911	255
9	31945	255
10	31933	255

o12.1

```

truncate table op_tch_results;
alter session set commit_wait=nowait;

declare
  loop_var      number;
  dbarfil_var    number := 6 ;
  dbablk_var     number := 2566541;
  tch_var        number ;
begin
  for loop_var in 1..600
  loop
    begin
      select tch
      into   tch_var
      from   x$bhh
      where  dbarfil = dbarfil_var
      and    dbablk  = dbablk_var;
      EXCEPTION WHEN OTHERS THEN tch_var := -1;
    end;

    insert into op_tch_results values ( loop_var, current_timestamp,
      dbarfil_var, dbablk_var, tch_var );

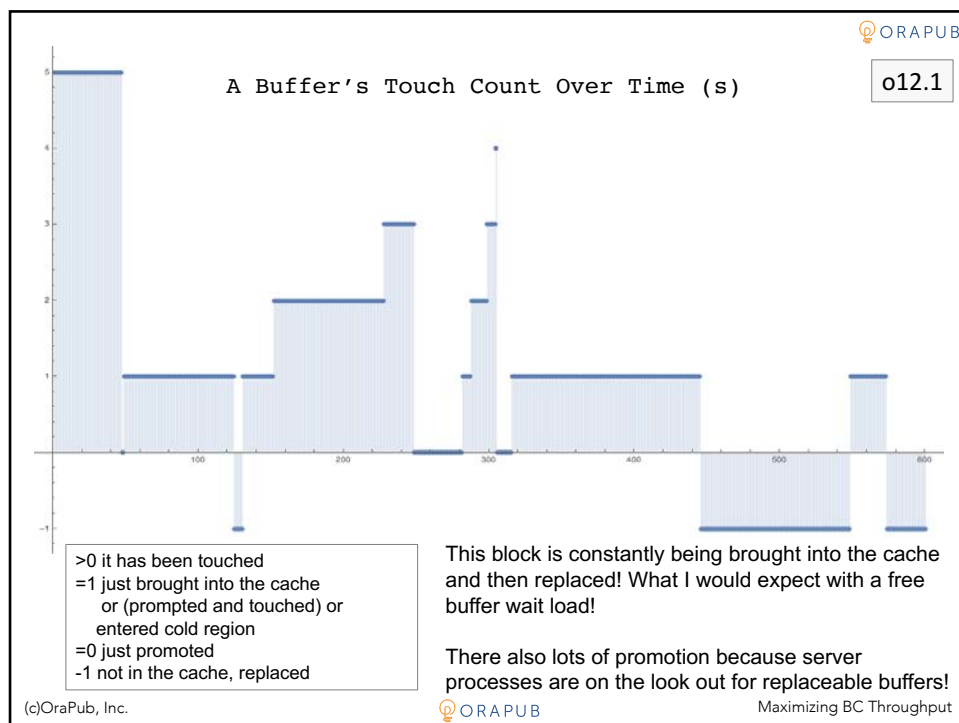
    dbms_lock.sleep(0.50);
  end loop;
end;
/


```

An active buffer in a FBW workload mix.

(c)OraPub, Inc.
ORAPUB
Maximizing BC Throughput

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.






This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

12c Touch Count Algorithm

- Based on what I previously showed:
 - Three second rule is still the default
 - Promotion still occurs when a server process or the DBWR is looking for a buffer AND the buffer's TC is ≥ 2 .
 - When promotion occurs, the TC is still reset to 0.
- Summary: operation remains the same.
- I did not test if the max 255 tch "rule" was in affect in 11g.

So, how long are the wait times... really


This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.



free buffer waits


This will happen if:

- All buffer gets have been suspended. This could happen when a file was read-only and is now read-write. All the existing buffers need to be invalidated since they are not linked to lock elements (needed when mounted parallel (shared)). So cache buffers are not assigned to data block addresses until the invalidation is finished.
- The session moved some dirty buffers to the dirty queue and now this dirty queue is full. The dirty queue needs to be written first. The session will wait on this event and try again to find a free buffer
- This also happens after inspecting **free buffer inspected** buffers. If no free buffer is found, Oracle waits for one second, and then tries to get the buffer again (depends on the context). For more information, see [free buffer inspected](#).

Wait Time: 1 second 

Parameter	Description
<i>file#</i>	See " file# "
<i>block#</i>	See " block# "

Oracle 12.1 Documentation

(c)OraPub, Inc.  Maximizing BC Throughput

```
SQL> @swEhVdGet 50000 10000 free%buffer%wait%

Collecting event wait times...

Copy and paste the below R code directly into the swEhCharts.r
It must be the last entry in the #2 area of swEhCharts.r

### START ###


eName = "free%buffer%wait%"

1 row selected.

### There were 50000 sampling loops completed and 1924 samples collected.
###
count0to1 = 3
count1to2 = 0
count2to4 = 2
count4to8 = 3
count8to16 = 1890
count16to32 = 0
count32to64 = 0
count64to128 = 0
count128to256 = 0
count256to512 = 0
count512to1024 = 0
count1024to2048 = 16
count2048to4096 = 11
count4096to8192 = 0
. . .

PL/SQL procedure successfully completed.

### END : Do NOT COPY and PASTE the "PL/SQL procedure..." TEXT INTO R ###
```



```
### START ###

eName = "free%buffer%wait%"

1 row selected.

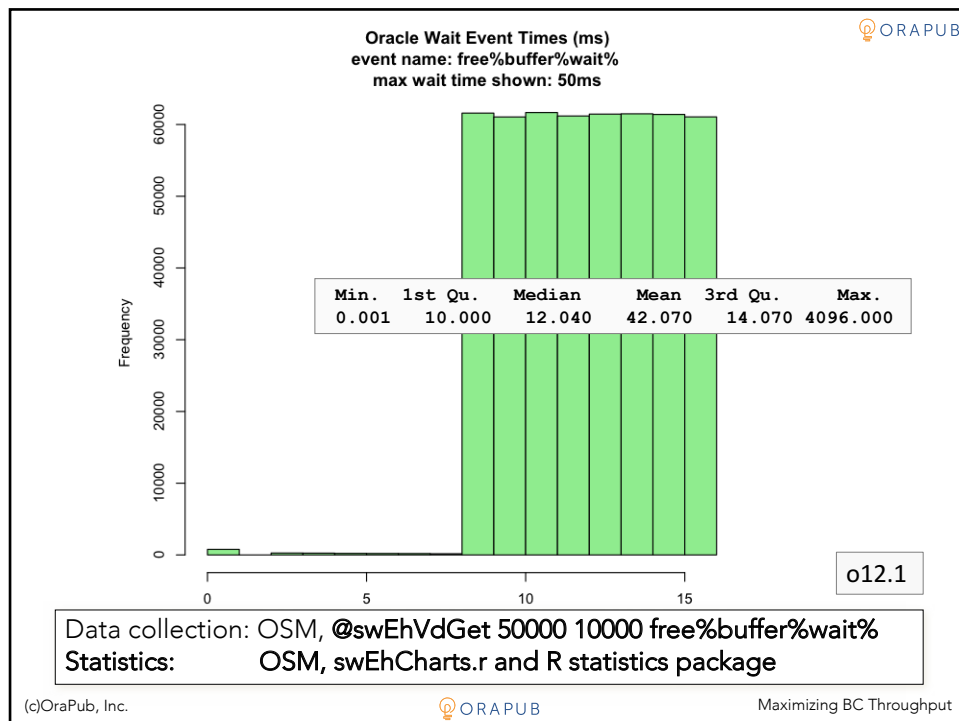
### There were 50000 sampling loops completed and 1924 samples collected.
###
count0to1 = 3
count1to2 = 0
count2to4 = 2
count4to8 = 3
count8to16 = 1890
count16to32 = 0
count32to64 = 0
count64to128 = 0
count128to256 = 0
count256to512 = 0
count512to1024 = 0
count1024to2048 = 16
count2048to4096 = 11
count4096to8192 = 0
. . .
```

(c)OraPub, Inc.



Maximizing BC Throughput

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.



How A Free Buffer Wait Occurs

1. Buffer replacement need – PIOR or CR
2. Shortage of non-popular free buffers - DML
3. DBWR falling behind

This presentation was given by Craig Shallahamer (craig@orapub.com)
at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

What is free buffer wait contention?



- An Oracle foreground process must find a replaceable buffer, but it is taking "too long." So, it gives up and yells, "Free Buffer Wait!" Waits and tries again.
- Remember, if the DBWR can't **pull** dirty buffers out of the buffer cache fast enough to make them free, then it is a **pull** from cache issue, not a **push** to disk issue.
- Don't focus on the IO subsystem first.

Watch how easy it is to cause FBWs



<https://youtu.be/pc2Wwi1cYAc>


(c)OraPub, Inc.



Maximizing BC Throughput

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Solutions
To Resolve
The Issue



(c)OraPub, Inc.




Maximizing BC Throughput

Resolving free buffer contention

- Add more DBWRs if available IO subsystem capacity. The DBWR can not get dirty buffers to disk fast enough, so increase DBWR *pull power* (e.g., more DBWRs) and perhaps some I/O write throughput improvement...
- Increase the buffer cache to increase the likelihood of finding a non-popular free buffer.
- Tune SQL to reduce PIO. By reducing PIO, you reduce the number of times a free buffer must be found.
- Find DML SQL to understand why there are so many dirty buffers. This may be an application issue or perhaps a workload balancing issue. If possible, reduce the number of changed buffers.
- There may be lots of clone (CR) buffer building, which requires a free buffer and also possibly additional free buffers for non-cached undo segment blocks. Look for SQL statements query a changing table, then tune the CR SQL to reduce CR buffers.
- Encourage the DBWR to write more often, but with smaller batches, by decreasing _db_large_dirty_queue.
- Be more patient, search longer, and give the DBWR more time to write its batch by increasing _db_writer_max_scan_pct.

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Solution Exploration: Increase Buffer Cache



Increasing the buffer cache

Report: ttpctx.sql

OSM by OraPub, Inc.

Page 1


Total Time Activity (88 sec interval)

Time Component	% TT	% WT	Avg Time Wait (ms)	Time (sec)	Wait Count(k)	
CPU consumption: Oracle SP + BG procs	46.63	0.00	0.000	272.881	0	
db file async I/O submit	12.58	23.57	343.925	73.600	0	
log file parallel write	11.99	22.47	76.941	70.170	1	
read by other session	8.39	15.71	0.575	49.070	85	
log buffer space	6.36	11.91	148.247	37.210	0	
db file scattered read	3.75	7.03	0.225	21.940	97	
control file parallel write	2.98	5.59	54.734	17.460	0	
log file switch complet				260	0	
log file switch (privat				040	0	
log file switch (check				540	0	
db file single write				990	0	
buffer busy waits	FBW %	20%	0%	810	2	
enq: CF - contention	BLK Chngs/s	19683	106838	230	0	
log file single write	LIO/s	150455	364635	960	0	
log file sync				490	0	
db file sequential read		0.01	0.02	0.006	0.060	10
PX qref latch		0.01	0.02	0.001	0.050	64

Statistic	Baseline	Incr BC	Incr DBWRs
FBW %	20%	0%	
BLK Chngs/s	19683	106838	
LIO/s	150455	364635	

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Solution
Exploration:
Add More
DBWRs




Add more DBWRs

Time Component	% TT	% WT	Avg Time Wait (ms)	Time (sec)	Wait Count(k)
db file async I/O submit	45.49	56.19	1647.284	382.170	0
CPU consumption: Oracle SP + BG procs	19.04	0.00	0.000	159.949	0
free buffer waits	16.43	20.29	16.314	138.000	8
log file parallel write	3.08	3.80	90.455	25.870	0
db file scattered read	2.33	2.88	0.017	19.610	1141
control file parallel				9.250	0
log buffer space				4.490	0
log file switch (priv)				3.290	0
log file switch (check)				2.260	0
db file single write	20%	0%	20%	1.950	0
buffer busy waits	BLK Chngs/s	19683	106838	23572	1.150
log file switch complete	LI0/s	150455	364635	165572	0.860
read by other session				0.860	16
enq: CF - contention	0.04	0.04	300.000	0.300	0
db file sequential read	0.03	0.04	0.003	0.260	89
log file single write	0.02	0.03	33.333	0.200	0
log file sync	0.02	0.02	85.000	0.170	0
Disk file operations I/O	0.01	0.01	7.143	0.050	0
cr request retry	0.01	0.01	0.006	0.050	8
PX qref latch					26


In this particular situation, there is a very good reason why adding DBWRs did not increase throughput.

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Moving To The Next Level



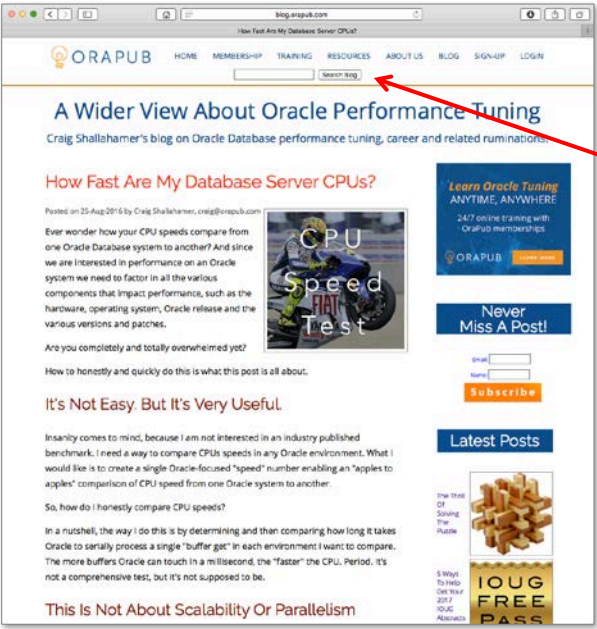
Download
the
FULL
presentation.




This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.


Resource Listing

- **OraPub Membership** for premium content
 - “How To” Webinars – two each month
 - Video Seminars – any device, any time, high quality
 - Community SLACK team
 - Learning paths, mentoring, assessments and certificates, priority response
- **Live Virtual Training – Multiple 2 hours sessions with daily break**
 - Oracle Tuning Fastpath
 - Oracle Buffer Cache Performance Analysis & Tuning
 - Tuning Oracle Using Advanced ASH Strategies
 - Tuning Oracle Using An AWR Report
 - Oracle Predictive Analysis Using Linear Regression
- **Craig’s Blog & Website** – Search: “lru”, “free buffer”
- **Toolkits** – OSM and BloodHound
- **Presentations** – Search OraPub.com: “lru”, “free buffer”, etc.
- **Books**
 - Oracle Performance Firefighting.
 - Forecasting Oracle Performance.

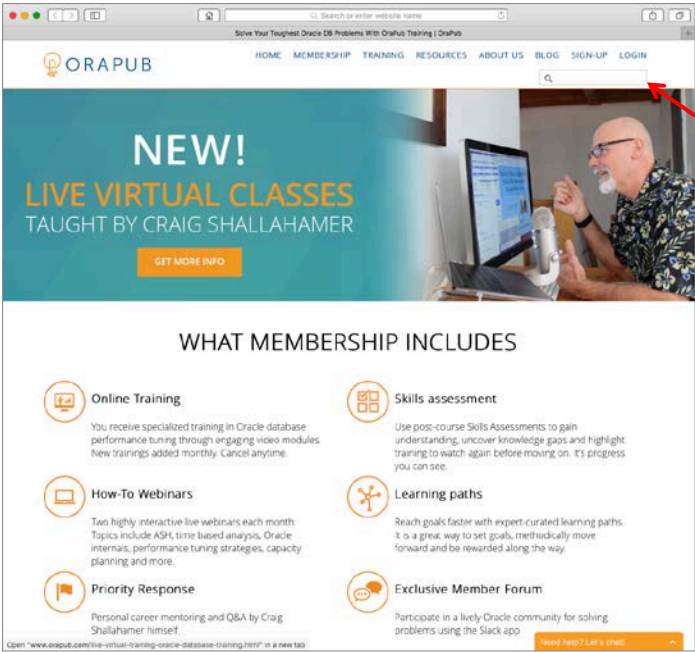





lru,
free
buffer,
visual


(c)OraPub, Inc.

Maximizing BC Throughput

This presentation was given by Craig Shallahamer (craig@orapub.com) at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.





lru,
free
buffer,
visual

(c)OraPub, Inc.

Maximizing BC Throughput

Thank You!

This presentation was given by Craig Shallahamer (craig@orapub.com)
at the November 2017 Bulgaria OUG (BGOUG) conference in Sophia, Bulgaria.

Maximizing **ORACLE®** CLOUD Buffer Cache Throughput

Craig Shallahamer
craig@orapub.com



ORACLE
ACE Director



ORAPUB

