

Top 10 Ways To Totally Mess Up An AWR Analysis

<https://www.orapub.com/ppts>

Craig Shallahamer
craig@orapub.com



COLLABORATE 19
TECHNOLOGY AND APPLICATIONS FORUM
FOR THE ORACLE COMMUNITY



REGISTER NOW AT COLLABORATE.IOUG.ORG



For the most updated version
of this presentation, go to:

www.orapub.com/ppts

ORAPUB

WORKLOAD REPOSITORY report for

DB Name	DB Id	Instance	Inst num	Startup Time	Release	RAC
PROD30	147730122	prod30	1	21-Sep-15 13:12	12.1.0.2	NO

Host Name	Platform	CPU's	Cores	Sockets	Memory (GB)
sixcore	Linux x86 64-bit	6	6	1	23.54

	Snap Id	Snap Time	Sessions	Cursors/Session
Begin Snap:	13009	21-Sep-15 13:41:28	45	2.2
End Snap:	13010	21-Sep-15 13:58:12	50	2.1
Elapsed:		16.72 (mins)		
DB Time:		400.28 (mins)		

Report Summary

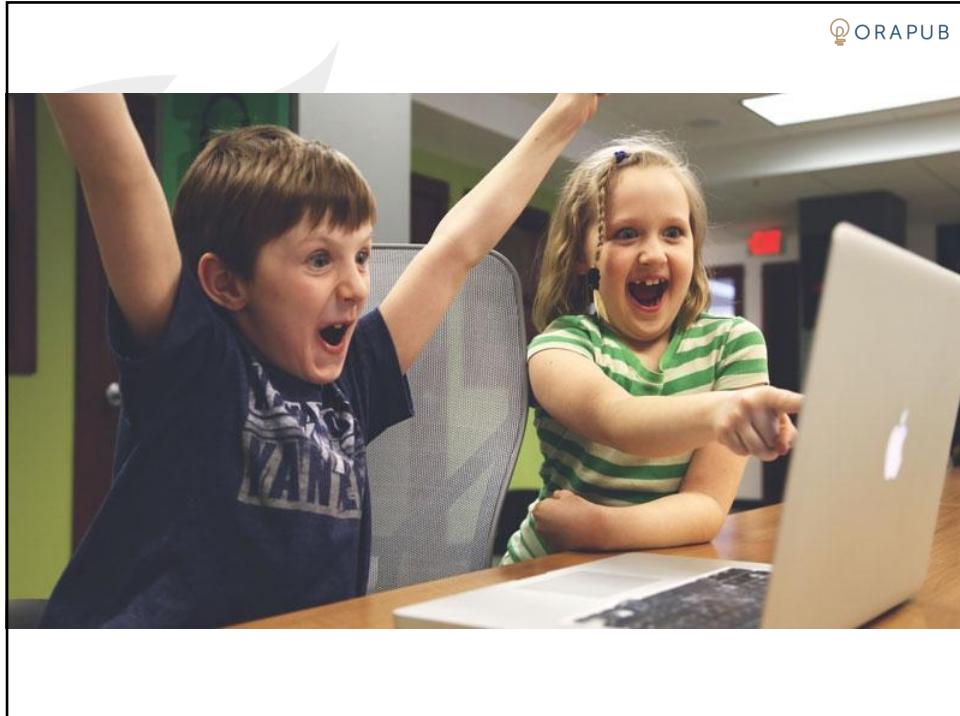
Cache Sizes

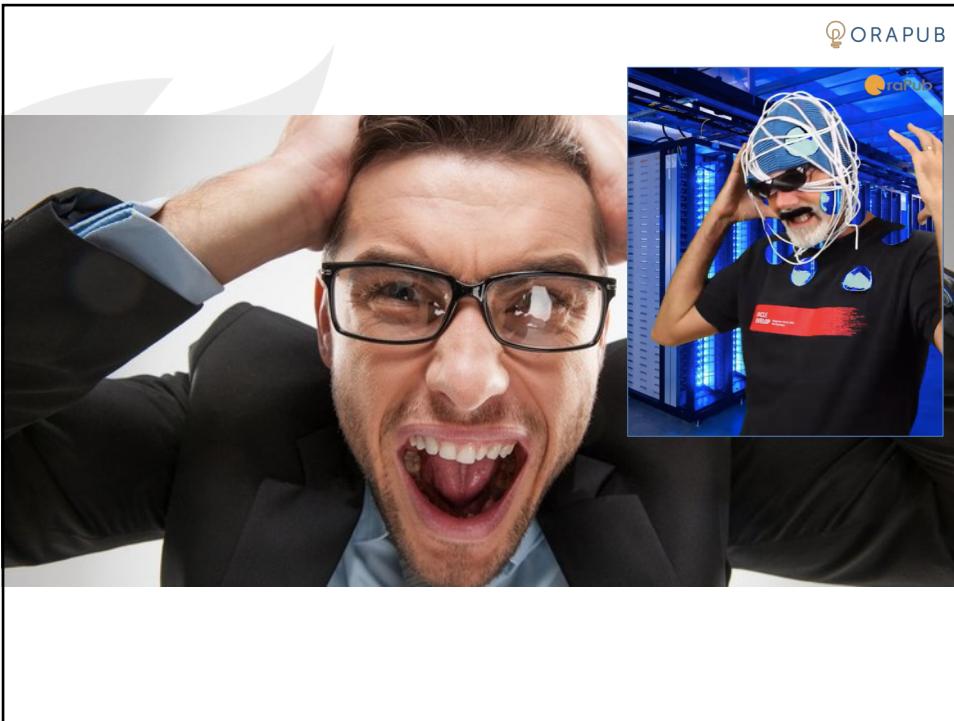
	Begin	End	
Buffer Cache:	0M	0M	Std Block Size: 8K
Shared Pool Size:	480M	480M	Log Buffer: 25,044K

Load Profile

	Per Second	Per Transaction	Per Exec	Per Call
DB Time(s):	23.9	255.5	0.00	30.83
DB CPU(s):	6.0	63.5	0.00	7.66
Redo size:	2,093.0	22,339.4		

(c)OraPub, Inc. AWR 10 Messups





By doing hundreds of
time based analysis,
I've noticed there are a
few key ways many
DBAs invalidate their
AWR analysis.

About Me...

- Long time Oracle DBA
- Specialize in Oracle Database performance and performance engineering
- Performance researcher
- Blogger: A Wider View About Oracle Performance Tuning
- Author: Oracle Performance Firefighting and Forecasting Oracle Performance.
- Conference speaker
- Teacher and mentor
- Oracle ACE Director
- IOUG DBA Track Manager

7

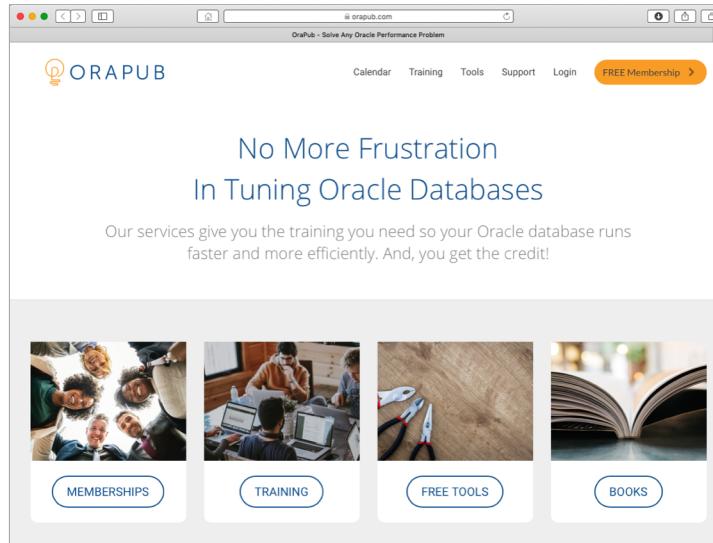
(c)OraPub, Inc.

ORAPUB



AWR 10 Messups

ORAPUB



The screenshot shows the OraPub website homepage. The header features the OraPub logo and navigation links for Calendar, Training, Tools, Support, and Login, along with a 'FREE Membership' button. The main headline reads 'No More Frustration In Tuning Oracle Databases'. Below the headline, a subtext states: 'Our services give you the training you need so your Oracle database runs faster and more efficiently. And, you get the credit!' A grid of four images with corresponding buttons: 'MEMBERSHIPS' (showing a group of people), 'TRAINING' (showing a person at a desk with laptops), 'FREE TOOLS' (showing a pair of pliers), and 'BOOKS' (showing an open book). The footer contains the number '8' and the text '(c)OraPub, Inc.'

OraPub works with Oracle DBAs empowering them to beat bots, AI, machine learning and autonomous anything.

AWR 10 Messups

It's all about time



because that's what
users care about

CPU Consumption Exclusion

1/10

Who gets the major award?

Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
DB CPU		47.7K		82.4	
cell smart table scan	1,440,243	4575	3.18	7.9	User I/O
cell single block physical read	3,741,750	2692.8	0.72	4.7	User I/O

Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% D
DB CPU		47.7K		
cell smart table scan	1,440,243	4575	3.18	
cell single block physical read	3,741,750	2692.8	0.72	
enq: KO - fast object checkpoint	5,746	1045.8	182.01	
log file sync	462,766	385.3	0.83	
reliable message	3,306	346	104.65	
flashback log file sync	1,077	143.5	133.22	
cell multiblock physical read	296,720	139.9	0.47	
direct path write	949	102.1	107.60	
cursor: pin S wait on X	19	62.2	3274.42	

11

(c)OraPub, Inc.

AWR 10 Messups

CPU Consumption Exclusion

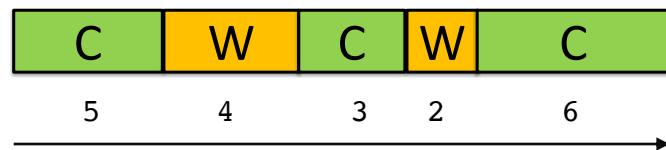
1/10

Elapsed Time Confusion

2/10

Oracle Time. What is it?

Database Time = CPU Time + Non-Idle Wait Time



T I M E

CPU = 14

Non-Idle Wait = 6

W O R K

300K LIOs

Elapsed = DB time = 20

Idle Wait = 0

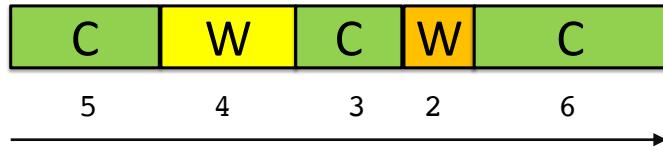
Experience = 20

Single Serial
Session

Oracle Time. What is it?



Database Time = CPU Time + Non-Idle Wait Time



TIME	WORK
CPU	= 14
Non-Idle Wait	= 2

Elapsed = DB time	= 16
Idle Wait	= 4
Experience	= 20

300K LIOs

Single Serial Session

Elapsed Time May Not Equal Wall Time

$$\text{Wall Time} = \frac{\text{Elapsed Time}}{\text{Parallelism}}$$

WORKLOAD REPOSITORY report for

DB Name	DB Id	Instance	Inst num	Startup Time	Release	RAC
BOX54	708114003	box54b		1 04-Feb-18 03:02	12.1.0.2.0	YES
Host Name	Platform	CPUs	Cores	Sockets	Memory (GB)	
sumertime001.westchck.edu	Linux x86 64-bit	72	36	2	755.34	
Snap Id	Snap Time	Sessions	Cursors/Session	Instances		
Begin Snap:	26871 07-Aug-18 08:00:06	2194	7.6	2		
End Snap:	26872 07-Aug-18 09:00:18	2251	8.0	2		
Elapsed:	60.20 (mins)					
DB Time:	963.42 (mins)					

This “Elapsed Time” is
NOT
 Oracle Time-Based “Elapsed Time”

SQL Statistics Yes!

- [SQL ordered by Elapsed Time](#)
- [SQL ordered by CPU Time](#)
- [SQL ordered by User I/O Wait Time](#)
- [SQL ordered by Gets](#)
- [SQL ordered by Reads](#)
- [SQL ordered by Physical Reads \(UnOptimized\)](#)
- [SQL ordered by Executions](#)
- [SQL ordered by Parse Calls](#)
- [SQL ordered by Sharable Memory](#)
- [SQL ordered by Version Count](#)
- [SQL ordered by Cluster Wait Time](#)
- [Complete List of SQL Text](#)

ORAPUB

SQL Statistics

- SQL ordered by Elapsed Time
- SQL ordered by CPU Time
- SQL ordered by User I/O Wait Time
- SQL ordered by Gets
- SQL ordered by Reads
- SQL ordered by Physical Reads (UnOptimized)
- SQL ordered by Executions
- SQL ordered by Parse Calls
- SQL ordered by Shareable Memory
- SQL ordered by Version Count
- SQL ordered by Cluster Wait Time
- Complete List of SQL Text

[Back to Top](#)

SQL ordered by Elapsed Time

Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.

Elapsed Time (s)	Executions	Elapsed Time per Exec (s)	%Total	%CPU	%IO	SQL Id	SQL Module	SQL Text
13,474.01	17,600	0.77	23.31	96.79	1.21	62r2rym6qb9x2	JDBC Thin Client	begin :1 := GRAIN_SOUP_CUST_FI...
12,385.62	3,201	3.87	21.43	99.71	0.03	dhan4yksznw4	JDBC Thin Client	SELECT DISTINCT HDOR.HDOR_ID A...
7,983.11	9,479	0.84	13.81	98.74	0.77	462ggd858h399	JDBC Thin Client	SELECT listagg(SOAP.SBSB_CK, ...
4,870.69	482	10.11	8.43	89.72	9.99	gzdmst9p1cf75	JDBC Thin Client	SELECT DISTINCT RSCL.*, TRIM(...
4,800.13	15,082	0.32	8.30	2.89	76.13	gdzc93uigmctp	JDBC Thin Client	SELECT TRANSACTION_ID, CREATED...
2,929.27	172	17.03	5.07	99.34	0.01	99sw69n1w3a24	JDBC Thin Client	SELECT PAYEE_TAX_ID, PAYEE_CHK...
2,137.05	1,375	1.55	3.70	98.00	0.25	6qk55fzvqmckr	JDBC Thin Client	SELECT listagg(SOAP.SBSB_CK, ...
1,642.74	4,269	0.38	2.84	96.48	0.48	2963trgg936vk	JDBC Thin Client	SELECT listagg(SOAP.SBSB_CK, ...
1,534.36	426	3.60	2.65	99.28	0.05	dd5778vkt0dy	JDBC Thin Client	SELECT DISTINCT RSCL.*, TRIM(...

Yes!

ORAPUB

SQL ordered by CPU Time

Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.

CPU Time (s)	Executions	CPU per Exec (s)	%Total	Elapsed Time (s)	%CPU	%IO	SQL Id	SQL Module	SQL Text
13,041.31	17,600	0.74	27.37	13,474.01	96.79	1.21	62r2rym6qb9x2	JDBC Thin Client	begin :1 := GRAIN_SOUP_CUST_FI...
12,349.30	3,201	3.86	25.92	12,385.62	99.71	0.03	dhan4yksznw4	JDBC Thin Client	SELECT DISTINCT HDOR.HDOR_ID A...
7,882.44	9,479	0.83	16.54	7,983.11	98.74	0.77	462ggd858h399	JDBC Thin Client	SELECT listagg(SOAP.SBSB_CK, ...
4,369.79	482	9.07	9.17	4,870.69	89.72	9.99	gzdmst9p1cf75	JDBC Thin Client	SELECT DISTINCT RSCL.*, TRIM(...
2,910.07	172	16.92	6.11	2,929.27	99.34	0.01	99sw69n1w3a24	JDBC Thin Client	SELECT PAYEE_TAX_ID, PAYEE_CHK...
2,094.20	1,375	1.52	4.39	2,137.05	98.00	0.25	6qk55fzvqmckr	JDBC Thin Client	SELECT listagg(SOAP.SBSB_CK, ...
1,584.98	4,269	0.37	3.33	1,642.74	96.48	0.48	2963trgg936vk	JDBC Thin Client	SELECT listagg(SOAP.SBSB_CK, ...

Yes!

Elapsed Time Confusion

2/10

Focusing on the insignificant

(looking for something interesting)

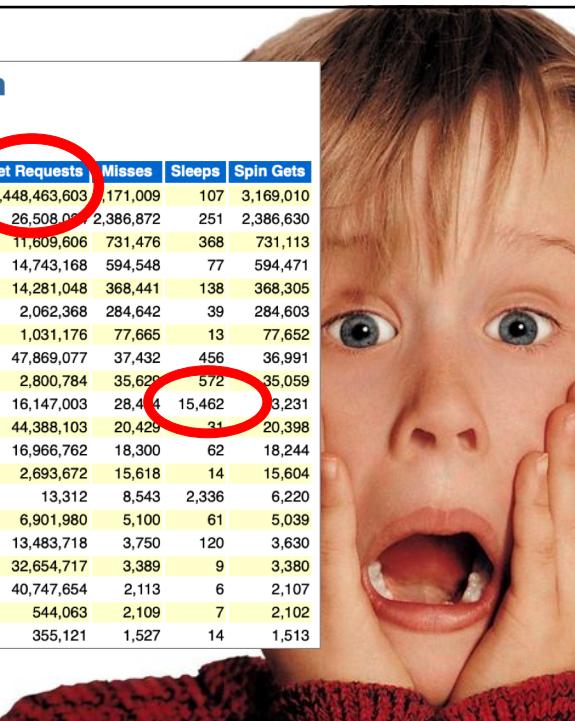
3/10

Latch Sleep Breakdown

- ordered by misses desc

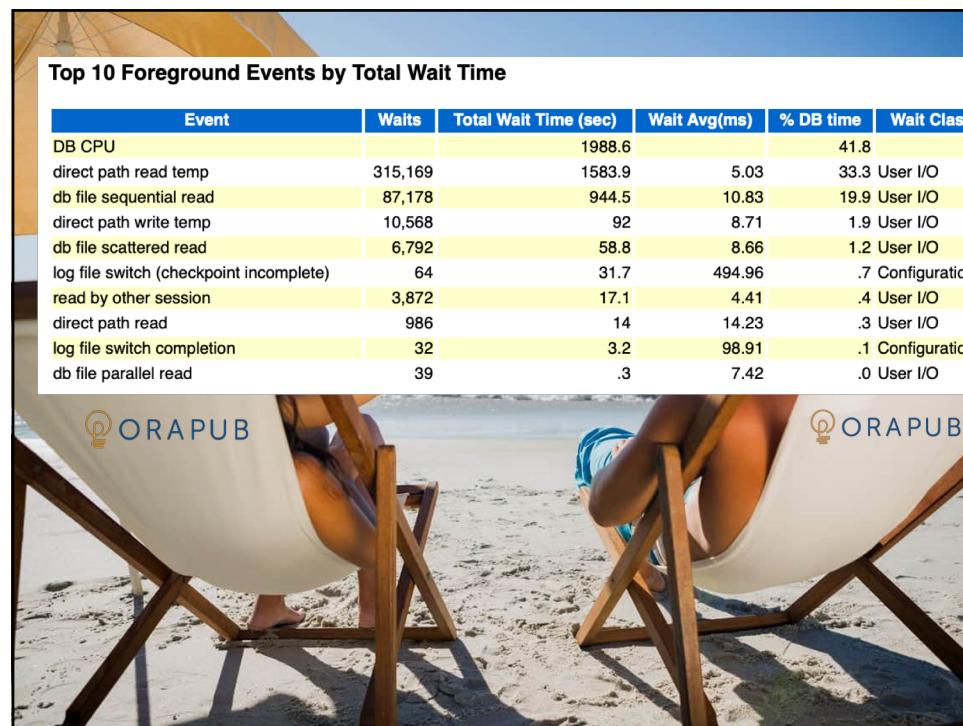
Latch Name	Get Requests	Misses	Sleeps	Spin Gets
cache buffers chains	8,448,463,603	1,171,009	107	3,169,010
messages	26,508,006	2,386,872	251	2,386,630
active checkpoint queue latch	11,609,606	731,476	368	731,113
enqueue hash chains	14,743,168	594,548	77	594,471
redo writing	14,281,048	368,441	138	368,305
LGWR NS Write	2,062,368	284,642	39	284,603
readredo stats and histogram	1,031,176	77,665	13	77,652
ges resource hash list	47,869,077	37,432	456	36,991
redo allocation	2,800,784	35,629	572	35,059
gc element	16,147,003	28,444	15,462	3,231
row cache objects	44,388,103	20,429	31	20,398
cache buffers lru chain	16,966,762	18,300	62	18,244
flashback sync request	2,693,672	15,618	14	15,604
space background task latch	13,312	8,543	2,336	6,220
ges cached resource lists	6,901,980	5,100	61	5,039
shared pool	13,483,718	3,750	120	3,630
ges enqueue table freelist	32,654,717	3,389	9	3,380
ASM map operation hash table	40,747,654	2,113	6	2,107
lgwr LWN SCN	544,063	2,109	7	2,102
KJC message pool free list	355,121	1,527	14	1,513

23



Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
DB CPU		1988.6		41.8	
direct path read temp	315,169	1583.9	5.03	33.3	User I/O
db file sequential read	87,178	944.5	10.83	19.9	User I/O
direct path write temp	10,568	92	8.71	1.9	User I/O
db file scattered read	6,792	58.8	8.66	1.2	User I/O
log file switch (checkpoint incomplete)	64	31.7	494.96	.7	Configuration
read by other session	3,872	17.1	4.41	.4	User I/O
direct path read	986	14	14.23	.3	User I/O
log file switch completion	32	3.2	98.91	.1	Configuration
db file parallel read	39	.3	7.42	.0	User I/O



Focusing on the insignificant

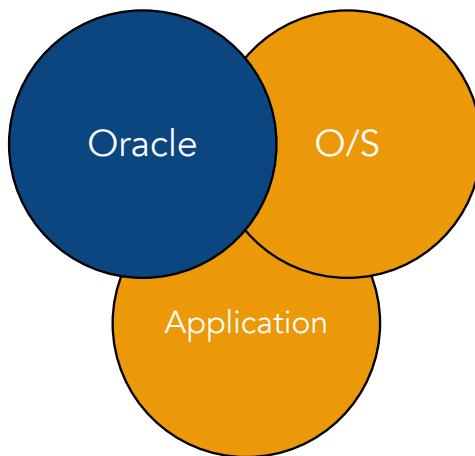
(looking for something interesting)

3/10

I'm A DBA
NOT
An OS Admin

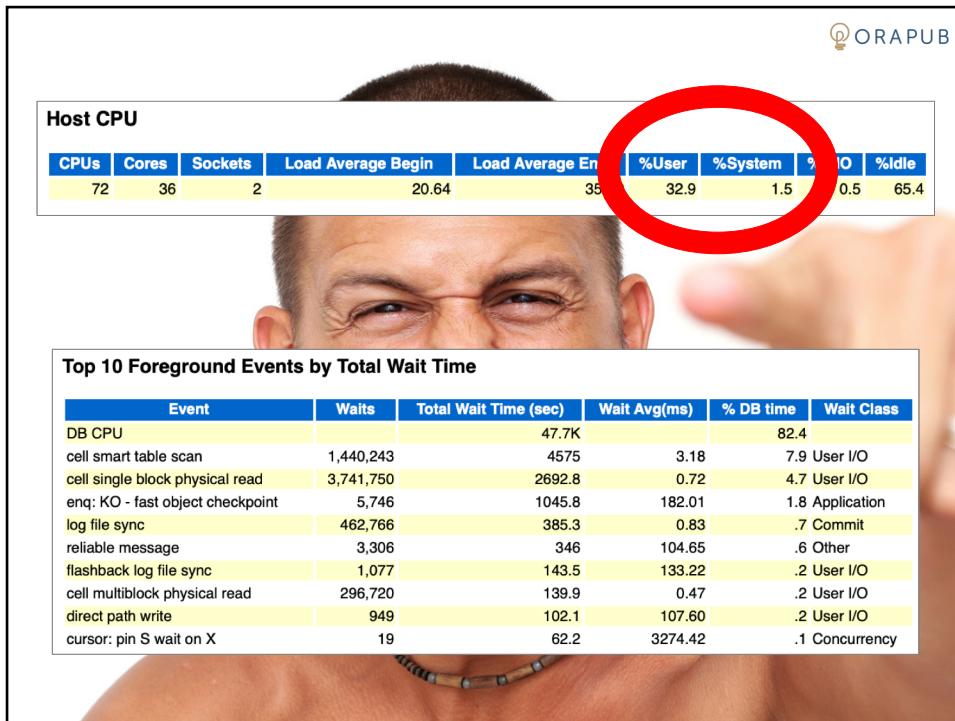
4/10

It's more than just Oracle.



Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
DB CPU		47.7K		82.4	
cell smart table scan	1,440,243	4575	3.18	7.9	User I/O
cell single block physical read	3,741,750	2692.8	0.72	4.7	User I/O
enq: KO - fast object checkpoint	5,746	1045.8	182.01	1.8	Application
log file sync	462,766	385.3	0.83	.7	Commit
reliable message	3,306	346	104.65	.6	Other
flashback log file sync	1,077	143.5	133.22	.2	User I/O
cell multiblock physical read	296,720	139.9	0.47	.2	User I/O
direct path write	949	102.1	107.60	.2	User I/O
cursor: pin S wait on X	19	62.2	3274.42	.1	Concurrency



Host CPU

CPUs	Cores	Sockets	Load Average Begin	Load Average End	%User	%System	%IO	%idle
72	36	2	20.64	35.0	32.9	1.5	0.5	65.4

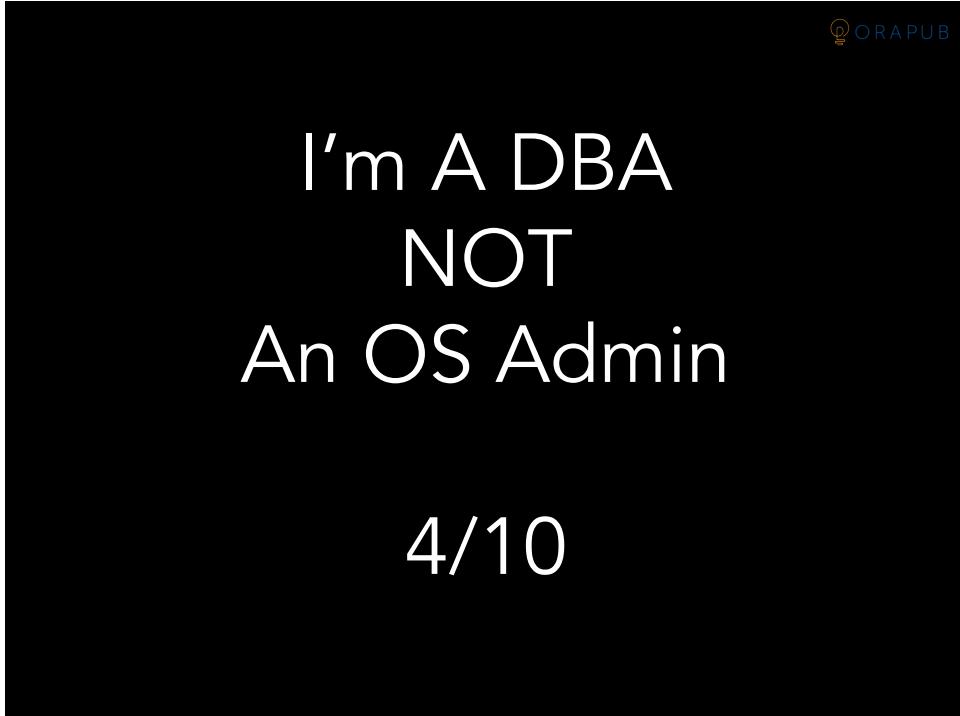
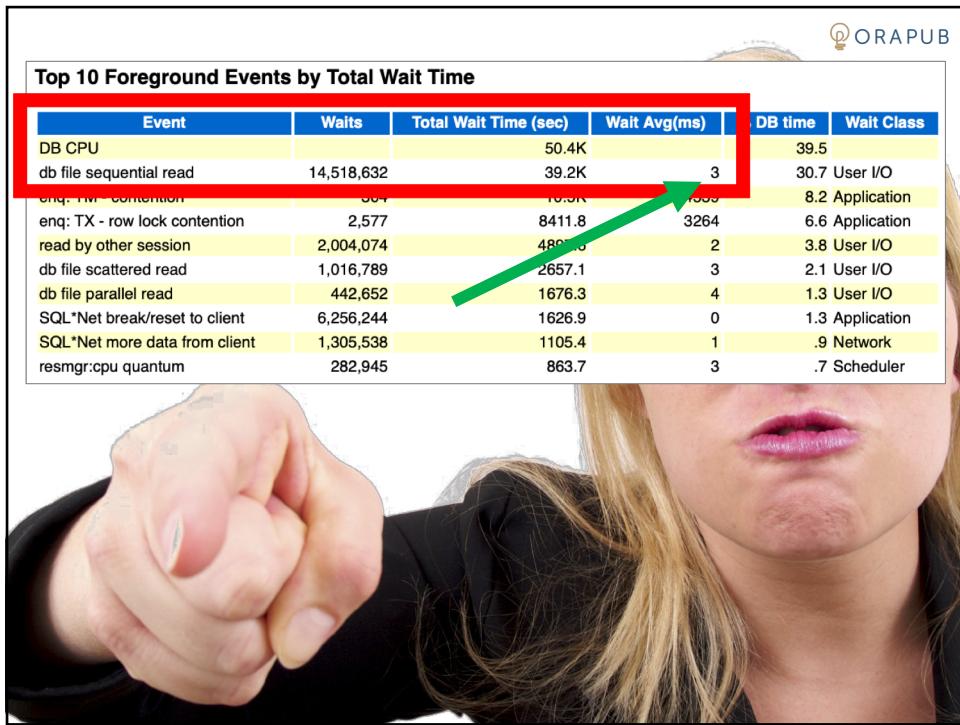
Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
DB CPU		47.7K		82.4	
cell smart table scan	1,440,243	4575	3.18	7.9	User I/O
cell single block physical read	3,741,750	2692.8	0.72	4.7	User I/O
enq: KO - fast object checkpoint	5,746	1045.8	182.01	1.8	Application
log file sync	462,766	385.3	0.83	.7	Commit
reliable message	3,306	346	104.65	.6	Other
flashback log file sync	1,077	143.5	133.22	.2	User I/O
cell multiblock physical read	296,720	139.9	0.47	.2	User I/O
direct path write	949	102.1	107.60	.2	User I/O
cursor: pin S wait on X	19	62.2	3274.42	.1	Concurrency



Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
DB CPU		50.4K		82.4	
db file sequential read	14,518,632	39.2K		7.9	User I/O
enq: TM - contention	337	10.3K		4.7	User I/O
enq: TX - row lock contention	2,577	8411.8		1.8	Application
read by other session	2,004,074	4897.6		.7	Commit
db file scattered read	1,016,789	2657.1		.6	Other
db file parallel read	442,652	1676.3		.2	User I/O
SQL*Net break/reset to client	6,256,244	1626.1		.2	User I/O
SQL*Net more data from client	1,305,538	1105.0		.1	Concurrency
resmgr:cpu quantum	282,945	863.0			



When an AWR report analysis is the wrong analysis strategy

5/10

Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
DB CPU		18.3K		25.0	
db file sequential read	595,835	18K	30	24.7	User I/O
log file sync	26,296	13.9K	530	19.1	Commit
SQL*Net message from dblink	3,230	7780.4	2409	10.7	Network
TCP Socket (KGAS)	122,399	7631.2	62	10.5	Network
Disk file operations I/O	76,981	3641.5	47	5.0	User I/O
direct path read	167,212	2316.2	14	3.2	User I/O
read by other session	26,883	1112.4	41	1.5	User I/O
buffer busy waits	3,499	711	203	1.0	Concurrency
direct path write temp	4,389	593.1	135	.8	User I/O

**By design,
an AWR report focuses on
an Oracle instance or cluster.**

34 (c)OraPub, Inc. AWR 10 Messups

AWR report says
CPU and physical
read IO is the heart
of the issue.

Users say
performance is
slow every Friday
afternoon.

AWR report says
CPU and physical
read IO is the heart
of the issue.

A few users said
their screen “just
froze” this
morning at 10:30.

Strategy? AWR or ASH or either

Characteristic	AWR	ASH
Repeatable?	Yes - Expected	No - Random
Problem Scope	Instance or RAC	Anything
Number of users	Must drive the "current"	Can be a single session, many sessions or all sessions
Need session details	No – Not possible	Yes – Session level samples
Need Wait Times	Yes – Average and more	No – Except 20+ sec

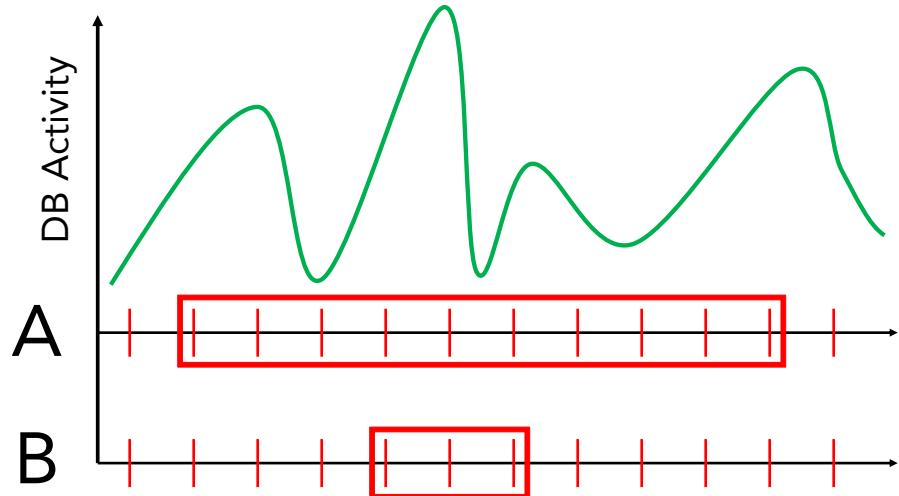
When an AWR report analysis is the wrong analysis strategy

5/10

Too BIG or SMALL snapshot interval

6/10

Which snapshot interval?



Which snapshot interval?



DB Activity
A

The general strategy is to set the begin and end snapshot IDs as closely as possible around the peak problem time.

B

The further before or after the snap IDs are set, the more “averaging” there will be, which can mask the issue.

B

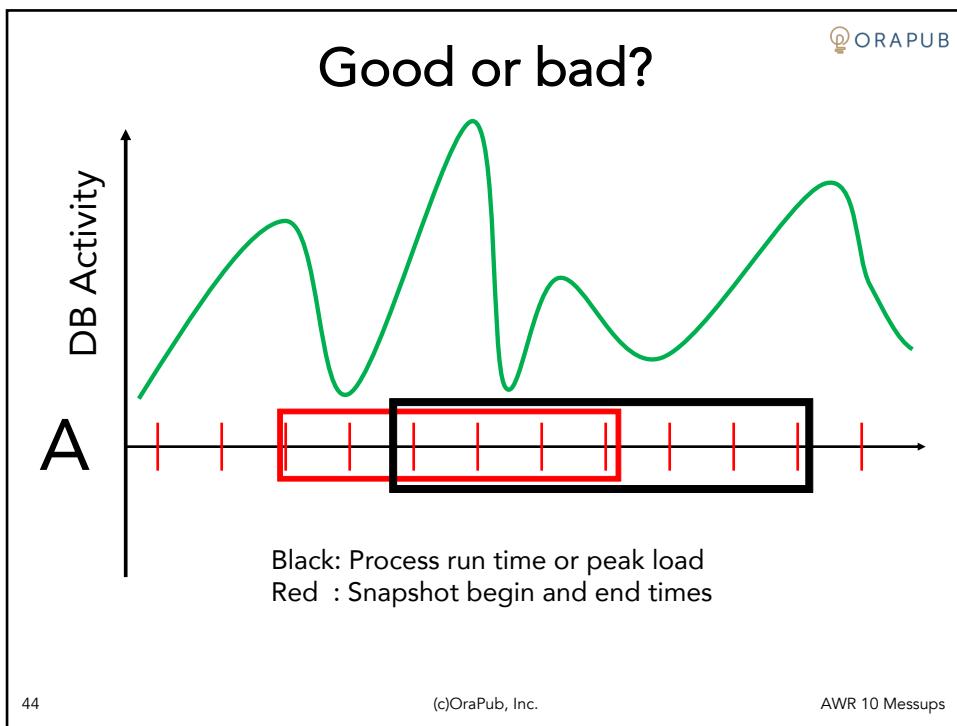
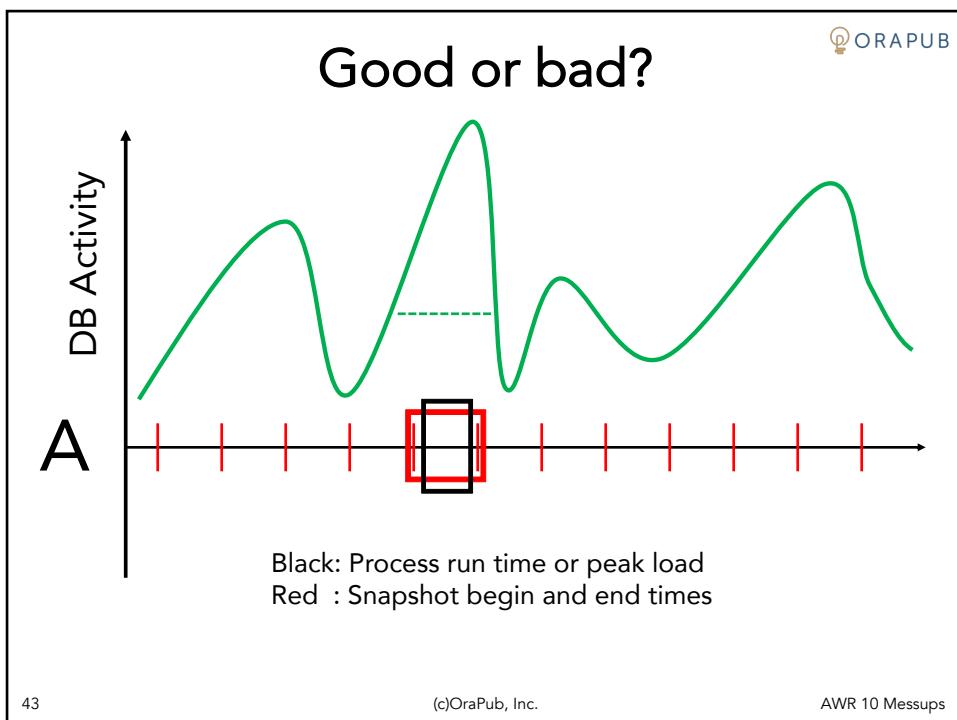
If we don't cover the “problem area” the results will not represent well the “problem area.”

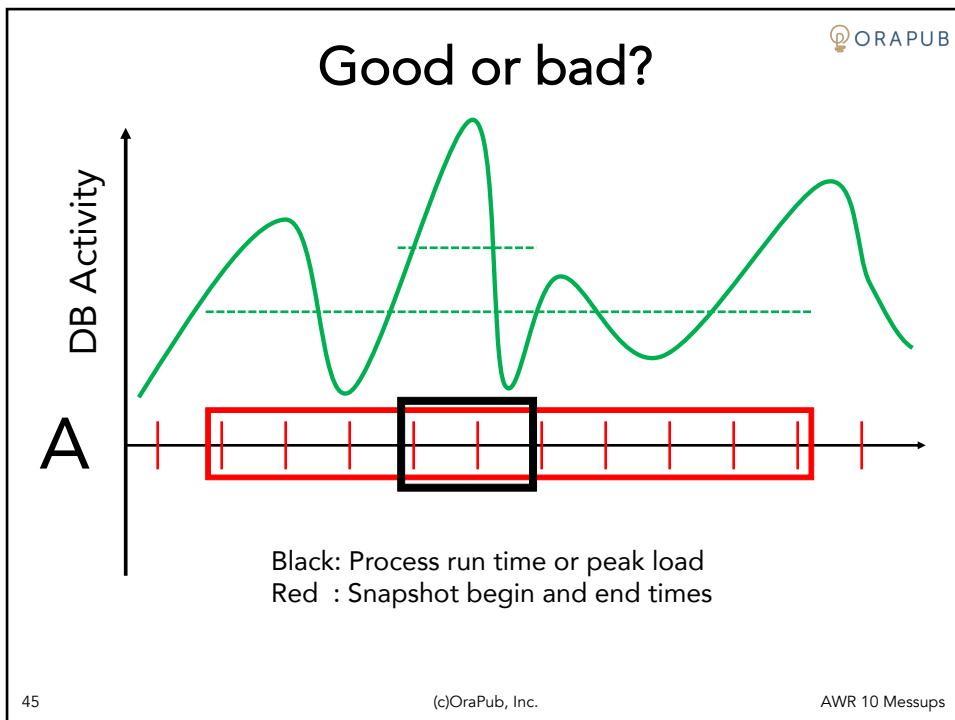
Good or bad?



DB Activity
A

Black: Process run time or peak load
Red : Snapshot begin and end times





Too
BIG or _{SMALL}
snapshot interval

6/10

Not Defining Top SQL

7/10

Just look at my
“TOP SQL”
monitor!



Or, use the column, **TOP_SQL_RANK**

```
SQL> desc dba_hist_sqlstat
Name Null? Type
-----
SNAP_ID NOT NULL NUMBER
DBID NOT NULL NUMBER
INSTANCE_NUMBER NOT NULL NUMBER
SQL_ID NOT NULL VARCHAR2(13)
PLAN_HASH_VALUE NOT NULL NUMBER
OPTIMIZER_COST NUMBER
OPTIMIZER_MODE VARCHAR2(10)
OPTIMIZER_ENV_HASH_VALUE NUMBER
SHARABLE_MEM NUMBER
TOP_SQL_RANK NUMBER
LOADED VERSIONS NUMBER
VERSION_COUNT NUMBER
MODULE VARCHAR2(64)
```

Or, use the column, **TOP_SQL_RANK**

```
SQL> desc dba_hist_sqlstat
Name Null? Type
-----
SNAP_ID NOT NULL NUMBER
DBID NOT NULL NUMBER
INSTANCE_NUMBER NOT NULL NUMBER
SQL_ID NOT NULL VARCHAR2(13)
PLAN_HASH_VALUE NOT NULL NUMBER
OPTIMIZER_COST NUMBER
OPTIMIZER_MODE VARCHAR2(10)
OPTIMIZER_ENV_HASH_VALUE NUMBER
SHARABLE_MEM NUMBER
TOP_SQL_RANK NUMBER
LOADED VERSIONS NUMBER
VERSION_COUNT NUMBER
MODULE VARCHAR2(64)
```

There is no AWR `TOP_SQL_RANK` column.

So, what does “top SQL” mean anyway?



Let your time-based analysis, the user experience and the OS situation drive and define the “top” SQL.

Top 10 Foreground Events by Total Wait Time

Event	Waits	Total Wait Time (sec)	Wait Avg(ms)	% DB time	Wait Class
DB CPU		47.7K		82.4	
cell smart table scan	1,440,243	4575	3.18	7.9	User I/O
cell single block physical read	3,741,750	2692.8	0.72	4.7	User I/O
enq: KO - fast object checkpoint	5,746	1045.8	182.01	1.8	Application
log file sync	462,766	385.3	0.83	.7	Commit
reliable message	3,306	346	104.65	.6	Other
flashback log file sync	1,077	143.5	133.22	.2	User I/O
cell multiblock physical read	296,720	139.9	0.47	.2	User I/O
direct path write	949	102.1	107.60	.2	User I/O
cursor: pin S wait on X	19	62.2	3274.42	.1	Concurrency

How should we define the “top SQL”?

DB Time is dominated by CPU consumption. So, CPU can be one way to define the “top SQL.”

SQL ordered by CPU Time

- Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.
- %Total - CPU Time as a percentage of Total DB CPU
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User IO Time as a percentage of Elapsed Time
- Captured SQL account for 88.0% of Total CPU Time (s): 47,652
- Captured PL/SQL account for 31.9% of Total CPU Time (s): 47,652

CPU Time (s)	Executions	CPU per Exec (s)	%Total	Elapsed Time (s)	%CPU	%IO	SQL Id	SQL Module	SQL Text
13,041.31	17,600	0.74	27.37	13,474.01	96.79	1.21	62r2rym6qb9x2	JDBC Thin Client	begin :1 := GRAIN_SOUP_CU
12,349.30	3,201	3.86	25.92	12,385.62	99.71	0.03	dhan4yksnnw4	JDBC Thin Client	SELECT DISTINCT HDOR.HI
7,882.44	9,479	0.83	16.54	7,983.11	98.74	0.77	462qgd858h389	JDBC Thin Client	SELECT listagg(SOAP.SBSB,
4,369.79	482	9.07	9.17	4,870.69	89.72	9.99	gzmst9p1cf75	JDBC Thin Client	SELECT DISTINCT RSCL.* ,
2,910.07	172	16.92	6.11	2,929.27	99.34	0.01	99sw69njw3a24	JDBC Thin Client	SELECT PAYEE_TAX_ID, PA

Could the top CPU SQL also be the top IO read and/or top Elapsed Time SQL?

SQL ordered by CPU Time

- Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.
- %Total - CPU Time as a percentage of Total DB CPU
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User IO Time as a percentage of Elapsed Time
- Captured SQL account for 88.0% of Total CPU Time (s): 47,652
- Captured PL/SQL account for 31.9% of Total CPU Time (s): 47,652

CPU Time (s)	Executions	CPU per Exec (s)	%Total	Elapsed Time (s)	%CPU	%IO	SQL Id	SQL Module	SQL Text
13,041.31	17,600	0.74	27.37	13,474.01	96.79	1.21	62r2rym6qb9x2	JDBC Thin Client	begin :1 := GRAIN_SOUP_CU
12,349.30	3,201	3.86	25.92	12,385.62	99.71	0.03	dhan4yksnnw4	JDBC Thin Client	SELECT DISTINCT HDOR.HI
7,882.44	9,479	0.83	16.54	7,983.11	98.74	0.77	462qgd858h389	JDBC Thin Client	SELECT listagg(SOAP.SBSB,
4,369.79	482	9.07	9.17	4,870.69	89.72	9.99	gzmst9p1cf75	JDBC Thin Client	SELECT DISTINCT RSCL.* ,
2,910.07	172	16.92	6.11	2,929.27	99.34	0.01	99sw69njw3a24	JDBC Thin Client	SELECT PAYEE_TAX_ID, PA

The top CPU SQL could also be the top elapsed time SQL and the top buffer get SQL.

SQL ordered by Elapsed Time

- Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.
- % Total - Elapsed Time as a percentage of Total DB time
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User I/O Time as a percentage of Elapsed Time
- Captured SQL account for 86.3% of Total DB Time (s): 57,805
- Captured PL/SQL account for 27.9% of Total DB Time (s): 57,805

Elapsed Time (s)	Executions	Elapsed Time per Exec (s)	%Total	%CPU	%IO	SQL Id	SQL Module	SQL Text
13,474.01	17,600	0.77	23.31	96.79	1.21	6212r7ym6qb9x2	JDBC Thin Client	begin :1 := GRAIN_SOUP_C...
12,385.62	3,201	3.87	21.43	99.71	0.03	dhan4yksznnw4	JDBC Thin Client	SELECT DISTINCT HDOR.H...
7,983.11	9,479	0.84	13.81	98.74	0.77	462qgd858h389	JDBC Thin Client	SELECT listagg(SOAP.SBS...
4,870.69	482	10.11	8.43	89.72	9.99	gzmst9p1cf75	JDBC Thin Client	SELECT DISTINCT RSCL.* ,
4,800.13	15,082	0.32	8.30	2.89	76.13	gdzc93ujgmctp	JDBC Thin Client	SELECT TRANSACTION_ID,...
2,929.27	172	17.03	5.07	99.34	0.01	99sw69njw3a24	JDBC Thin Client	SELECT PAYEE_TAX_ID, PA...

Just look at my
“TOP SQL”
monitor!

Stop
And
Think



Not Defining Top SQL

7/10

Not Using Parallelism To Your Advantage

8/10

Q: Which full bucket will empty first?



5 Gallons

5 Gallons



59

(c)OraPub, Inc.

AWR 10 Messups

Q: Which query will complete first?

5M LIOs

5M LIOs

No PQ



Yes PQ



60

(c)OraPub, Inc.

AWR 10 Messups

Make sure
your IO
subsystem can
absorb the
increased IO
rate, but for a
shorter period
of time.



Which SQL will likely benefit from increased parallelism?

SQL ordered by Elapsed Time

- Resources reported for PL/SQL code includes the resources used by all SQL statements called by the code.
- % Total DB Time is the Elapsed Time of the SQL statement divided into the Total Database Time multiplied by 100
- %Total - Elapsed Time as a percentage of Total DB time
- %CPU - CPU Time as a percentage of Elapsed Time
- %IO - User I/O Time as a percentage of Elapsed Time
- Captured SQL account for 86.3% of Total DB Time (s): 57,805
- Captured PL/SQL account for 27.9% of Total DB Time (s): 57,805

Elapsed Time (s)	Executions	Elapsed Time per Exec (s)	%Total	%CPU	%IO	SQL Id	SQL Module	SQL Text
13,474.01	17,600	0.77	23.31	96.79	1.21	62r2rym6qb9x2	JDBC Thin Client	begin :1 := GRAIN_SOUP_CI
12,385.62	3,201	3.87	21.43	99.71	0.03	dhan4yksznw4	JDBC Thin Client	SELECT DISTINCT HDR.H
7,983.11	9,479	0.84	13.81	98.74	0.77	462qgd858h389	JDBC Thin Client	SELECT listagg(SOAP.SSB
4,870.69	482	10.11	8.43	89.72	9.99	gzmst9p1cf75	JDBC Thin Client	SELECT DISTINCT RSCL.* ,
4,800.13	15,082	0.32	8.30	2.89	76.13	gdzc93ujgmctp	JDBC Thin Client	SELECT TRANSACTION_ID
2,929.27	172	17.03	5.07	99.34	0.01	99sw69njw3a24	JDBC Thin Client	SELECT PAYEE_TAX_ID, PA



Not Using Parallelism To Your Advantage

8/10

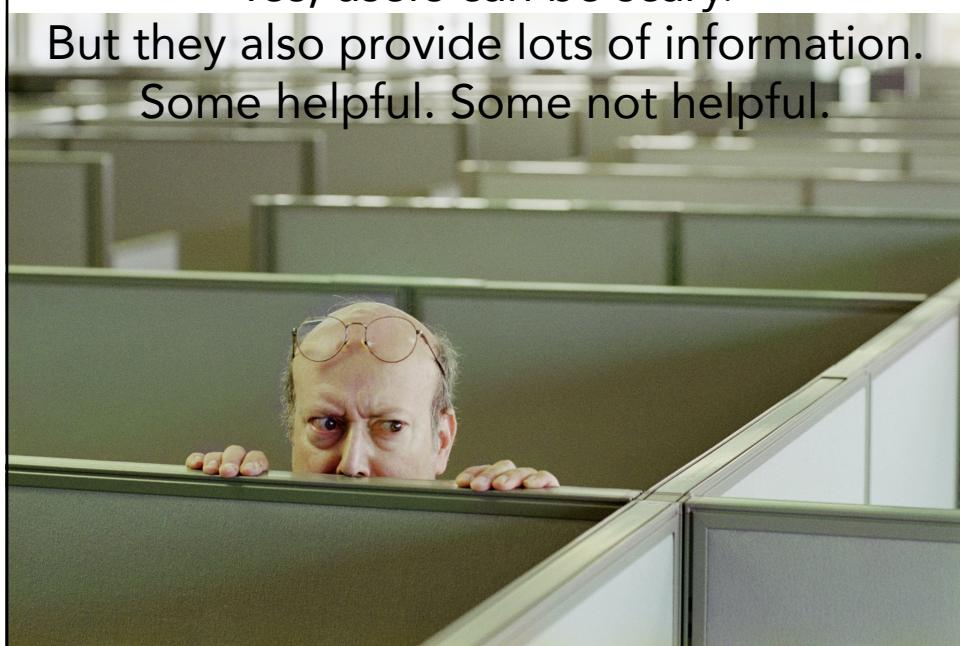
Avoiding The User Community

9/10

Yes, users can be scary.
But they also provide lots of information.
Some helpful. Some not helpful.



Yes, users can be scary.
But they also provide lots of information.
Some helpful. Some not helpful.

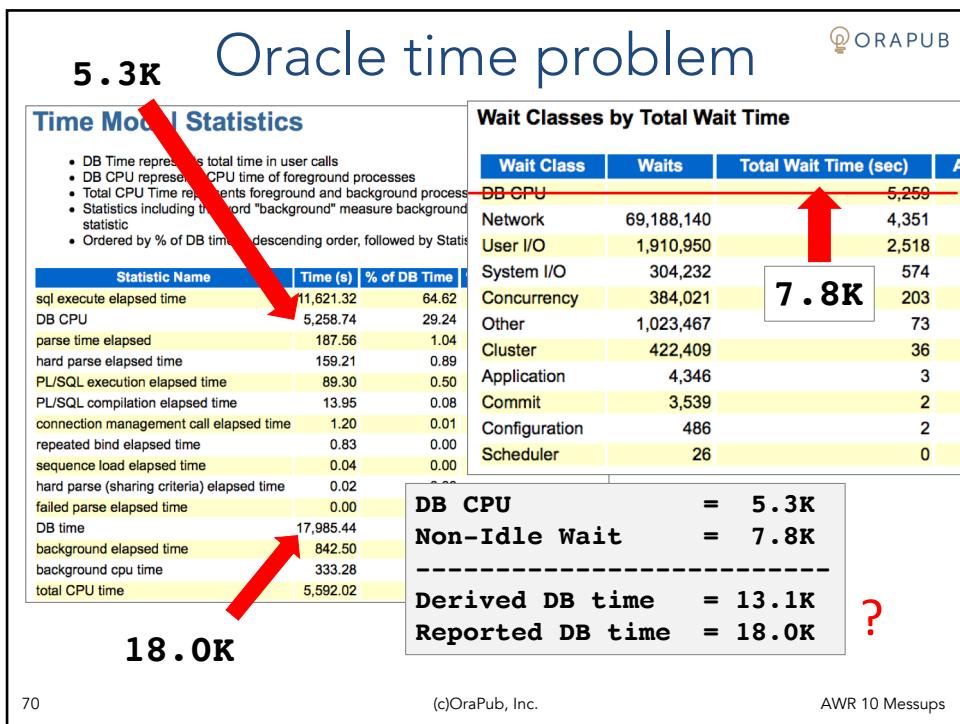
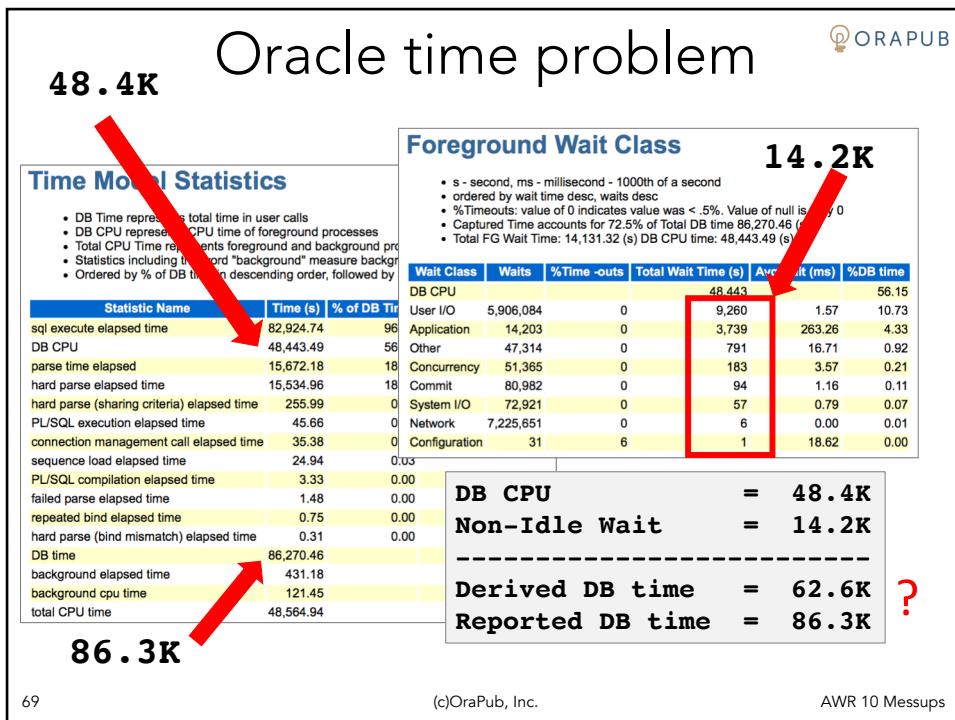


Avoiding The User Community

9/10

Blindly Trust Oracle Time Data

10/10



Do the math. If it's doesn't match...

- Use ASH. Because ASH is sample based and not time based, it avoids many of the time related problems.
- Use Operating System statistics (dba_hist_osstat) to validate DB CPU.
- When the OS is bottlenecked, strange times can be recorded. Just querying and recording the AWR data, can take awhile, impacting the data quality.

Blindly Trust Oracle Time Data

10/10



Download ALL my presentations:
www.orapub.com/ppts

www.orapub.com/free-webinars

There are FREE
AWR webinars.

Tuning Oracle Using An AWR Report



OraPub LVCs Are Different

You learn to master the topic at a deeper, higher confidence and more practical level compared to any other teaching method I offer.

This occurs because the class is spread out over multiple weeks, each session is 2 hours with a day in between, you do homework/activation on your real systems and I personally work with you through the entire class.

Details & Registration: <https://www.orapub.com/lvc>
Event Calendar: <https://www.orapub.com/events>

Thank You!



Top 10 Ways To Totally Mess Up An AWR Analysis

<https://www.orapub.com/ppts>

Craig Shallahamer
craig@orapub.com



COLLABORATE 19

TECHNOLOGY AND APPLICATIONS FORUM
FOR THE ORACLE COMMUNITY



APRIL 7-11, 2019

SAN ANTONIO

REGISTER NOW AT COLLABORATE.IOUG.ORG



For the most updated version
of this presentation, go to:

www.orapub.com/ppts